

Effectiveness of adversarial inputs on class-imbalanced convolutional neural networks

Rafael Possas, Ying Zhou

School of Information Technologies, University Of Sydney, Camperdown, NSW
{rafael.possas,ying.zhou}@sydney.edu.au
<https://sydney.edu.au>

Abstract. Convolutional neural networks (CNNs) performance has increased considerably in the last couple of years. However, as with most machine learning methods, these networks suffer from the data imbalance problem - when the underlying training dataset is comprised of an unequal number of samples for each label/class. Such imbalance enforces a phenomena known as domain shift that causes the model to have poor generalisation when presented with previously unseen data. Recent research has focused on a technique called *gradient sign* that intensifies domain shift in CNNs by modifying inputs to deliberately yield erroneous model outputs, while appearing unmodified to human observers. Several commercial systems rely on image recognition techniques to perform well. This wide use, therefore, has heightened the need for better understanding of this threat. In this work we present an experimental study that sheds light on the link between adversarial attacks, imbalanced learning and transfer learning. Through a series of experiments we evaluate the fast gradient sign method on class imbalanced CNNs, linking model vulnerabilities to the characteristics of its underlying training set and internal model knowledge.

Keywords: convolutional neural networks, adversarial examples, gradient sign, imbalanced training, transfer learning

1 Introduction

Convolutional neural networks (CNNs) are a class of non-linear machine learning algorithms known for its state of the art performance on datasets with spatial structure. To date, not much research has been done into the adversarial inputs on CNNs - a process on which inputs are changed to manipulate the algorithm outputs. The motivation for adversarial robustness comes largely from being able to shield image recognition systems from behaving unexpectedly. Experimental demonstrations of the effectiveness of adversarial attacks were carried out mainly by [2], [5], [16] and have highlighted the need for improvement on the current state of CNNs techniques. Developing robustness to such attacks has become of the utmost importance as many commercial applications are based on the same small group of models. However, these previously published studies are limited to

show the general effectiveness of adversarial methods rather than understanding the deep relationship with the underlying training set distribution.

Domain shift or dataset shift [17] is also a well known cause for low performance of several machine learning algorithms [7], [8]. This happens when the joint distribution of inputs and outputs differs between training and test stages, causing models to perform badly on unseen data. The adverse effect of domain shift is even worse on real world as data distributions is often skewed and rarely contains enough information to learn all the required features of the data domain. Adversaries have been proven to more readily exploit domain shift [12], [10], the question as to whether imbalanced training sets affects adversarial inputs performance on CNNs is still unanswered.

The effectiveness of an adversarial attack is also dependent on the internal gradient information of the targeted model. As shown on Papernot et al (2016), attacks could be classified as both black-box and white-box. The former uses gradient information from a separate model, while the latter uses the target model gradient to generate adversarial inputs. While black-box attack is approximation of the internal information of the target, the white box uses the true representation of that model feature space.

Currently, there is no empirical evidence on the effectiveness of adversarial inputs on class-imbalanced CNNs. We designed a set of experiments to investigate the effects of both training sets with skewed distributions and the model's internal gradient information on the robustness of these networks to such attacks. The main contributions of this work are as follows:

1. Evaluation of the resilience of imbalanced CNNs to adversarial attacks using the white-box and black box approach.
2. Investigation of classes with overlapping distributions and their relationship to both adversarial attacks and the imbalanced learning problem.

Section 2 of this paper discusses the related work in both CNNs, gradient sign methods, adversarial attacks and imbalanced/transfer learning. Section 3 provides details of the training models, imbalanced datasets and gradient sign methods used in our experiments. Section 4 presents the results on the under-sampled, over-sampled and balanced cases using both same/different model gradient. Sections 5 is dedicated to drawing conclusions and providing directions to related future work.

2 Related work

Previous work has shown that the high-dimensional non-linearities of convolutional neural networks [11] creates adversarial pockets of space - places where datapoints could be moved in order to provide a wrong model output . As shown on Figure 1, such pockets enables methods to deliberately create an adversary that produces an incorrect, high confidence prediction for an image without visible distortion [13]. This can be done by adding intentional noise to each pixel of




Fig. 1. Adversarial exploration of pockets of spaces [14]

an image so as to fool the algorithm into predicting an incorrect label [5], [15], [19].

The gradient sign method was introduced by Goodfellow et al. (2014) and has been used as the foundation of many of the experiments in adversarial attacks on CNNs. The results have shown that convolutional neural networks have linear behavior in very high dimensional spaces [5]. Most inputs were miss-classified not only by Goodfellow et. al (2014) experiments but others as well [2],[16].

The work of Papernot et al. (2016) has shown that one can use transfer learning to perform white and black box attacks against CNNs [15], [20] and, thus, to intentionally force the model to predict specific labels. The combination of adversaries and transfer learning creates a threat vector for many state of the art methods. Attacks can, however, depend on some specific internal information of the target model [12], [15]. As most recent applied methods depend on the network gradient information, there is a strong dependence on the network confidence per class label and the robustness of the model to adversarial attacks.

Techniques to overcome imbalanced learning have been developed for more general machine learning models. The work of Heibo et al [6], for instance, provides a technique for doing weighted sampling of minority classes to minimize the effect of imbalanced learning. Another approach could be to incorporate unsupervised clustering on synthetic data generation mechanism in order to avoid wrong generation of synthetic samples [1]. More recent work has used a Bayesian framework to increase l_2 robustness to adversarial examples [2].

3 Experiment design

Our experiments aims to investigate the relationships of the underlying learning structure of CNNs and the perturbation caused by gradient sign methods. In particular we focus on the investigation of how the gradient step from the sign method moves the points away from their distributions, and how this could be affected by both balanced and imbalanced distributions. This requires class labels of the data set to be non-hierarchical so we can make better assumptions of their distributions.

We use the CIFAR-10 data set [9] in our experiment. CIFAR-10 data is visually rich and empowers the analysis between different class labels. The data set contains 32x32 images in 10 classes, each has 5,000 samples for training and 1,000 for testing. There is not much overlap nor hierarchical relationship between classes. The 2014 ImageNet dataset [4] would be the natural choice. However, its hierarchically organized categories adds unnecessary complexity to the experiment design and makes it hard to establish the causality relationship.

3.1 Network architecture and synthetic dataset imbalance

Network Architecture - All the experiments were done using a modified VGGNet architecture [18]. The two fully connected 4096 layers at the end were replaced by one single layer with 512 neurons and RELU activations. In addition, the total number of convolutions blocks and pooling layers were reduced to 3, with the first layer having 2 stacked convolution layers followed by a max pooling of stride 2x2 and the last two layers with 3 stacked convolutions also followed by a max pooling of stride 2x2. We have used RMSProp [3] as the optimisation technique with a learning rate of 10^{-4} and a decay 10^{-5} . Figure 2 shows that our model has an overall accuracy of approximately 83%, which is comparable to many state of the art models nowadays.




Fig. 2. Results of our adapted VGG architecture on the CIFAR-10 dataset, shows comparable overall performance

Dataset imbalancing As the CIFAR-10 dataset is not naturally imbalanced, we have artificially created two variations on which we trained the imbalanced networks. One dataset consists of a direct under-sample of the target class to

1,000 samples, and the other was changed using an oversampling of the target class (or an under-sampling of all other classes). We kept the number of samples for the target class at 5,000 while all other classes were reduced to 1,000 samples. For each class of the two different datasets configurations, a network was then trained until convergence using the same hyper-parameters as the balanced case. Each model was evaluated against a test set of 1,000 samples of the target class which was perturbed by its own under/over-sampled model and the balanced model. The two sources of gradient information are referred as white-box and black-box attacks since the former has complete information of the network weights and biases while the other uses an approximation of the same parameters.

Both imbalanced networks were separately tested for each class on white-box and black-box adversarial attacks. The white-box test was designed to investigate the vulnerability of class imbalance on adversarial examples while the black-box test is designed to verify the robustness on transfer learning environments. In total we evaluated 50 different combinations: 20 for each different imbalanced dataset (same model gradient and balanced network gradient) and 10 for the balanced network using its own gradients on each class. Figure 3 shows the accuracy for the models without any perturbation. It can be seen that the individual class accuracy for the under-sampled case is reduce while the same metric is increased on the over-sampling network.

3.2 Gradient sign methods

The gradient sign is a method that uses internal gradient information to create directed perturbation to input data. The resulting a label will be different whether one adds or subtracts noise according to equations 1 and 2.

$$C(x + \delta) \approx C(x) + \epsilon * \text{sign}(\nabla C) \quad (1)$$

$$C(x + \delta) \approx C(x) - \epsilon * \text{sign}(\nabla C) \quad (2)$$

The gradient sign equation has a simple interpretation. The main goal is to add a change δ into each pixel of the image so as to make that image look like more or less likely to the chosen gradient extracted from the source network. The sign on our ∇C indicates that we are only interested on the direction of the gradient while the ϵ controls the magnitude of the step.

Suppose the current true label of the class is selected as a gradient candidate, adding noise would mean that we increase the cost function of our input while subtracting noise is the same as minimizing our loss function even further. The equations above are usually referred as ascent and descent methods

Perturbations could also be applied by two variations of the gradient sign method. While the fast gradient sign method applies a single perturbation to the input, the iterative gradient sign method performs the same perturbation a chosen number of times iteratively [5]. Figure 4 shows an example of adversarial created using the fast method.

In order to enforce consistency throughout our experiments, we have chosen the true sample label as the backpropagated gradient along with the fast gradient




Fig. 3. Individual class accuracy for under-sampled, over-sampled case on the CIFAR-10 modified dataset shows a decrease in accuracy for classes with lower number of samples




Fig. 4. Adversarial example crafting with fast gradient sign [5].

sign ascent method. The intuition behind this choice is that we look to increase the cost function of the target class by moving away from the current true label. The ϵ value chosen was 0.01 as it provided the best trade-off between misclassification rate and the amount of visible change applied to the input image.

4 Results

We use the results of the balanced network on adversarial attacks as the baseline to evaluate whether imbalanced CNNs are less or more vulnerable to adversarial learning. Table 1 shows that the accuracy for all classes is drastically reduced when the balanced model is presented with adversarial examples. Models with under-sampled datasets were even more vulnerable than balanced networks. Figure 5 shows the relative difference for all the three different networks (balanced, under-sampled and over-sampled). Values were calculated by finding the difference between the perturbed accuracy and the non-perturbed accuracy of each class model. They represent the percentage on which the initial accuracy was reduced. The under-sampled model had the higher relative difference on average, which shows that the imbalanced nature of the dataset ended-up increasing the vulnerability of the model.

Perturbation on the over-sampling case had a weaker effect, as the small push caused by our ϵ was not enough to move points to outside of their distributions. Objects of the over-sampled classes would need bigger steps in order




Fig. 5. Relative difference for each model. Higher numbers means more vulnerability

to successfully create an adversarial that leads to a wrong classification label. Accuracy for most of the over-sampling cases was around 45% and the relative difference was the lowest of all three models, which shows robustness of the target over-sampled class. Class imbalanced models are naturally affected by the

Class Label	Black-box		White-box		
	Undersample	Oversample	Balanced	Undersample	Oversample
0 - Airplane	60%	87%	36%	19%	61%
1 - Automobile	64%	91%	23%	16%	63%
2 - Bird	38%	73%	20%	9.4%	27%
3 - Cat	21%	72%	11%	0.5%	19%
4 - Deer	58%	80%	20%	9.8%	20%
5 - Dog	47%	76%	15%	9%	38%
6 - Frog	76%	88%	27%	20%	49%
7 - Horse	59%	88%	20%	18%	52%
8 - Ship	69%	89%	37%	19%	59%
9 - Truck	46%	87%	49%	21%	54%

Table 1. Results for the two different sources of perturbations along with the two different imbalanced datasets. Under-sampling intensifies adversarial attack while over-sampling increases model robustness

false positive and false negative trade off shown on figure 6. The decision boundaries on such models favour the class with more samples and, hence, increases the accuracy for one class while decreasing for the other classes. The area under the curve for misclassified examples on the under-sampled distribution is bigger, and it is caused by the suboptimal exploration of feature space of that class. This effect is exploited by adversaries as there is an increase on the misclassification rate of distributions with lower amplitude.

The increased number of samples of the over-sampled label causes the network to perform a trade-off when optimizing its loss function. For instance, the

decision boundary would be chosen in order to minimize the total error of the network. The cost function is lower when the decision boundary minimizes the misclassification of the majority class as there is a higher number of samples. The choice of a biased decision boundary could be one of the factors explaining the higher resilience of over-sampled networks.

4.1 Transfer learning and overlapping distributions

Transfer learning - The use of a different model gradient for creating adversaries has shown less effective when compared to the same model attack. As the overall gradient have not only different direction but also magnitudes, the system has proven to be more robust to the attack. The experiment reveals that although gradient sign is quite effective for fooling networks it does require a good amount of knowledge from the underlying training parameters so as to unleash its full potential. Attacking an under-sampled/over-sampled network with the gradient of the balanced network did not show to be as effective as using the same model's gradient. The average accuracy of an under-sampled model attack with adversaries generated from a different network was 53.8% while the same metric was 25.8% for the same model attack. Even that our training samples are within the same data domain, there are still huge differences on the gradients learned from the network.

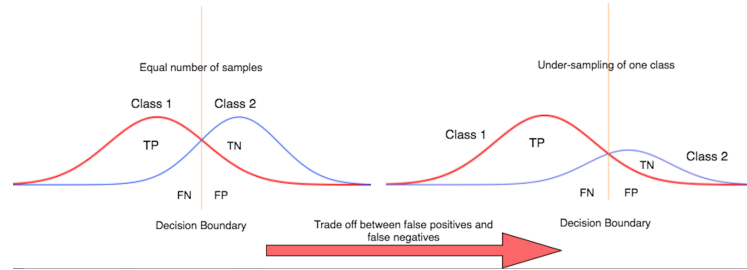


Fig. 6. Dataset imbalance causes models to perform adjustments of decision boundaries leading to an increase on accuracy of the majority class and decrease on the minority class.

Overlapping distributions - The results for class distributions with similar set of features show strong impact of the adversarial attack. Figure 2 shows that for the pairs cat/dog and automobile/truck there is already a natural misclassification between one another. For instance 13% of dog samples as misclassified as cat in the original balanced model. Our experiment shows that the adversarial attacks intensify this phenomena by increasing the number of times on which one class is picked over the other. Figure 7 shows that cats are increasingly misclassified as dogs when under-sampled dataset on the cat class is used. While on

the cat under-sampling case 40% of the samples were misclassified as dogs, on the oversampling one, 39% of the dogs were misclassified as cats. This outcome provides interesting insights, as it shows that the gradient sign is behaving linearly in the high dimensional space and 'moving' in the direction of the closest vicinity.



Fig. 7. From left to right: cat under-sampling / over-sampling with perturbation.

5 Conclusion and future work

We have shown that adversarial attacks are even more severe on datasets with under-sampled class labels and that the decision boundary trade-off on the over-sampled classes increases their robustness to adversarial examples. Labels with similar features have also shown higher vulnerability to the fast gradient sign methods as their similarities in the high dimensional space facilitates the technique to successfully create an adversary of the similar class.

As several commercial applications rely on almost the same group of models, understanding of such properties is of extreme importance. Future work in this field could look further in datasets with a higher number of classes and more complex relationships between labels so as to not only confirm our insights but also discover new interesting properties of CNNs.

References

1. Barua, S., Islam, M.M., Murase, K.: A Novel Synthetic Minority Oversampling Technique for Imbalanced Data Set Learning, pp. 735–744. Springer Berlin Heidelberg, Berlin, Heidelberg (2011), http://dx.doi.org/10.1007/978-3-642-24958-7_85

2. Billoovits, C., Eric, M., Agarwala, N.: Hitting depth: Investigating robustness to adversarial examples in deep convolutional neural networks (2016)
3. Dauphin, Y., de Vries, H., Bengio, Y.: Equilibrated adaptive learning rates for non-convex optimization. In: *Advances in Neural Information Processing Systems*. pp. 1504–1512 (2015)
4. Deng, J., Dong, W., Socher, R., Li, L.J., Li, K., Fei-Fei, L.: Imagenet: A large-scale hierarchical image database. In: *Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE Conference on*. pp. 248–255. IEEE (2009)
5. Goodfellow, I.J., Shlens, J., Szegedy, C.: Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572* (2014)
6. He, H., Bai, Y., Garcia, E.A., Li, S.: Adasyn: Adaptive synthetic sampling approach for imbalanced learning. In: *Neural Networks, 2008. IJCNN 2008. (IEEE World Congress on Computational Intelligence). IEEE International Joint Conference on*. pp. 1322–1328. IEEE (2008)
7. Japkowicz, N., Stephen, S.: The class imbalance problem: A systematic study. *Intelligent data analysis* 6(5), 429–449 (2002)
8. Krawczyk, B.: Learning from imbalanced data: open challenges and future directions. *Progress in Artificial Intelligence* 5(4), 221–232 (2016)
9. Krizhevsky, A.: Cifar-10 and cifar-100 datasets (2009), <https://www.cs.toronto.edu/~kriz/cifar.html>
10. Laskov, P., Lippmann, R.: Machine learning in adversarial environments. *Machine Learning* 81(2), 115–119 (2010), <http://dx.doi.org/10.1007/s10994-010-5207-6>
11. Lawrence, S., Giles, C.L., Tsoi, A.C., Back, A.D.: Face recognition: A convolutional neural-network approach. *IEEE transactions on neural networks* 8(1), 98–113 (1997)
12. Lowd, D., Meek, C.: Adversarial learning. In: *Proceedings of the Eleventh ACM SIGKDD International Conference on Knowledge Discovery in Data Mining*. pp. 641–647. KDD '05, ACM, New York, NY, USA (2005), <http://doi.acm.org/10.1145/1081870.1081950>
13. Papernot, N.: On the integrity of deep learning systems in adversarial settings (2016)
14. Papernot, N.: Machine learning in adversarial settings (2017), <https://www.papernot.fr/files/16-mcdaniel-sp-machine-learning-in-adversarial-settings.pdf>
15. Papernot, N., McDaniel, P., Goodfellow, I.: Transferability in machine learning: from phenomena to black-box attacks using adversarial samples. *arXiv preprint arXiv:1605.07277* (2016)
16. Papernot, N., McDaniel, P., Goodfellow, I., Jha, S., Berkay Celik, Z., Swami, A.: Practical black-box attacks against deep learning systems using adversarial examples. *arXiv preprint arXiv:1602.02697* (2016)
17. Quionero-Candela, J., Sugiyama, M., Schwaighofer, A., Lawrence, N.D.: *Dataset Shift in Machine Learning*. The MIT Press (2009)
18. Simonyan, K., Zisserman, A.: Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556* (2014)
19. Szegedy, C., Zaremba, W., Sutskever, I., Bruna, J., Erhan, D., Goodfellow, I., Fergus, R.: Intriguing properties of neural networks. *arXiv preprint arXiv:1312.6199* (2013)
20. Yosinski, J., Clune, J., Bengio, Y., Lipson, H.: How transferable are features in deep neural networks? In: *Advances in neural information processing systems*. pp. 3320–3328 (2014)